

Introduction to Distributed Systems

Vania Marangozova-Martin

Vania.Marangozova-Martin@imag.fr

Web Site: <http://ids.forge.imag.fr>

Agenda

Week	Tuesday, 09:45 – 11:15 (1h30)	Tuesday, 11:15 – 12:45	
5	01-févr	Introduction to distributed systems. Java RMI.	RMI Tutorial
		Rattrapage Tuesday 17h: Distributed Chat with RMI	
6	08-févr	Distributed Web applications. Servlets.	Servlet Tutorial.
		Rattrapage Tuesday 17h: Distributed chat RMI + demos	
7	15-févr	Servlets Lab	Servlets Lab + demos
8	22-févr	Interruption week	
9	29-févr	Coordination-Based Systems	Coordination Tutorial.
10	07-mars	No IDS (Vania away)	
11	14-mars	Coordination Lab	Coordination Lab + demos
12	21-mars	P2P System	P2P Tutorial.
13	28-mars	P2P Lab	P2P Lab
14	04-avr	Cloud Computing	Amazon, Google, Azure Cloud services presentations from students
15	11-avr	Cloud Lab (Part I)	Cloud Lab (Part II)
16	18-avr	Interruption week	

Goals

- ◆ Get an insight on the complexity of distributed systems
- ◆ Introduction to middleware
 - ❖ Why do we need it?
 - ❖ What is inside?
- ◆ Practical work
 - ❖ This course is "practice-oriented"
 - ❖ Learn to use modern technologies: web servers, noSQL databases, cloud...

Administrativa

- ◆ Evaluation
 - ❖ RMI chat demo
 - ❖ Servlets web application demo
 - ❖ Cloud demo
 - ❖ Presentation of a cloud service
 - ❖ Final exam
- ◆ Web site: ids.forge.imag.fr
- ◆ Do not hesitate to mail me questions
 - ❖ tagged in the subject with [M1_MOSIG_IDS]

What is a distributed system

- ◆ “A distributed system is one in which the failure of a computer you didn't even know existed can render your own computer unusable.”

Leslie Lamport, 1987.

Distributed Systems Characteristics


- ◆ Definition of a distributed system
 - ❖ Multiple components interconnected with a communication system; components have computing (CPU), storage (memory, disks) or interface functions (sensors)
 - ❖ The components are not independent but collaborate towards a common goal

Why distribution?

Distribution is everywhere!

- ◆ Systems need it
 - ❖ Integration of initially separated modules
 - ❖ Massive resource integration
 - Grids, clouds, data centers
 - ❖ New application domains
 - Ubiquitous computing
 - Surveillance, domotics
- ◆ Technical advances
 - ❖ Cost and performances
 - ❖ Generalized interconnection
 - Computers+TV+telecom
 - Sensor networks

What is the difference between a centralized and a distributed system? 1/7



centralized system

- ◆ All is accessible on the same machine
 - ❖ Memory
 - ❖ Data on disk
 - ❖ I/O
- ◆ Programs run locally
 - ❖ Can be monitored
 - Memory state
 - Process state
 - ...

What is the difference between a centralized and a distributed system? 2/7

Moniteur d'activité (Toutes les opérations)

Nom de l'opération	Mémoire	Mém. compressée	Fichs	Ports	PID	Utilisateur
kernel_task	1,39 Go	0 octets	163	0	0	root
Firefox	952,8 Mo	198,1 Mo	62	399	21092	vania
WindowServer	576,9 Mo	326,3 Mo	6	1 101	175	_windowserver
Photos	511,7 Mo	0 octets	5	341	21873	vania
Finder	343,5 Mo	281,8 Mo	22	951	259	vania
Mail	232,7 Mo	77,0 Mo	20	556	16919	vania
Microsoft PowerPoint	202,7 Mo	31,6 Mo	9	223	21692	vania
installld	199,1 Mo	198,6 Mo	2	68	847	root
photolibraryd	184,4 Mo	65,0 Mo	2	113	355	vania
Skim	157,5 Mo	123,0 Mo	11	330	15031	vania
JabRef	157,2 Mo	127,1 Mo	50	562	14865	vania
softwareupdated	152,3 Mo	147,2 Mo	12	209	235	_softwareupda
mds_store	145,8 Mo	88,8 Mo	7	71	190	root
iTunes	134,8 Mo	0 octets	21	422	21748	vania
Microsoft Excel	117,3 Mo	23,4 Mo	15	205	21700	vania
Aquamacs	113,4 Mo	59,4 Mo	5	231	20092	vania
com.apple.MediaLibraryService	109,7 Mo	0 octets	2	77	21765	vania
Calendar	101,2 Mo	46,1 Mo	3	214	14649	vania
iconservicesagent	93,1 Mo	90,6 Mo	5	88	288	vania
Dock	88,7 Mo	45,6 Mo	6	338	256	vania
https://su.itunes.apple.com	86,0 Mo	66,9 Mo	21	329	20577	vania
recentd	84,5 Mo	50,5 Mo	3	86	385	vania
Numbers	69,0 Mo	51,9 Mo	5	211	21480	vania
CalendarAgent	68,8 Mo	32,0 Mo	4	333	276	vania
mds	63,9 Mo	33,0 Mo	7	1 525	58	root
Agent Photos	61,4 Mo	21,4 Mo	8	225	10316	vania
ContentStoreMail	60,6 Mo	24,4 Mo	14	506	14895	vania

PRESSION SUR LA MÉMOIRE


Mémoire physique :	16,00 Go
Mémoire utilisée :	7,10 Go
Cache :	2,35 Go
Fichier d'échange utilisé :	1,72 Go

Mémoire de l'application : 4,46 Go
Mémoire résidente : 1,76 Go
Comprimé : 694,8 Mo

What is the difference between a centralized and a distributed system? 3/7

What can be distributed?

- The resources are distributed
 - CPU
 - Memory (live, persistent) is distributed and not shared



What is the difference between a centralized and a distributed system? 4/7

What can be distributed?

- The code




a variable here

is not accessible here

what we do here

is not what we do here



What is the difference between a centralized and a distributed system? 5/7

What can be distributed?


- The data



data is not here

data is not here


data is not here



What is the difference between a centralized and a distributed system? 6/7


♦ What can be distributed?

- ❖ Interactions



What is the difference between a centralized and a distributed system? 7/7

- ♦ No global time
- ♦ No global state
- ♦ ... and it is supposed to work anyway 😊

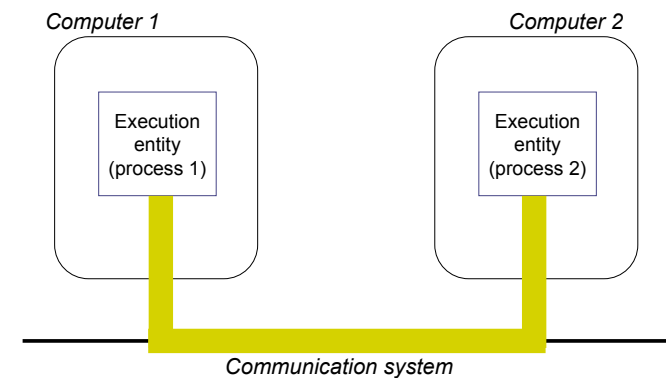


Distributed Systems Characteristics (2)

♦ Desired properties

- ❖ The system should be able to operate (at least in degraded mode) in the case of component failure
- ❖ The system should support communication failures (message loss, disconnection, ...)
- ❖ The system should resist to security attacks (confidentiality violations, integrity violations, denial of service, resource stealing, ...)

A Distributed System

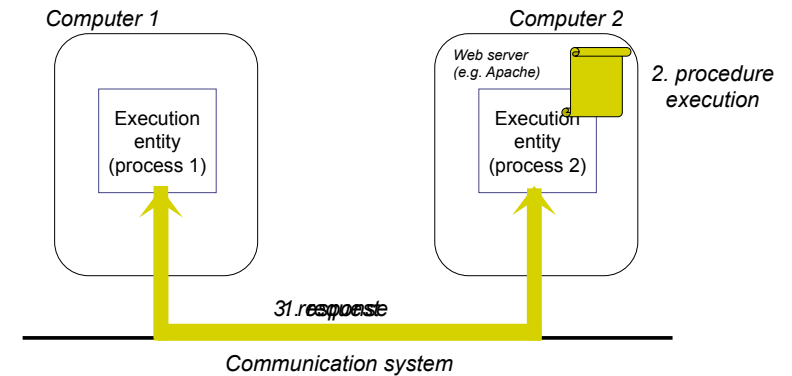


Communication mechanisms in a distributed system

- ◆ Direct (i.e. Synchronous) communication
 - ❖ Program to program
 - E.g. remote procedure call
 - ❖ Program to database
 - E.g. distributed transaction processing
- ◆ Indirect (i.e. Asynchronous) communication
 - ❖ Message passing

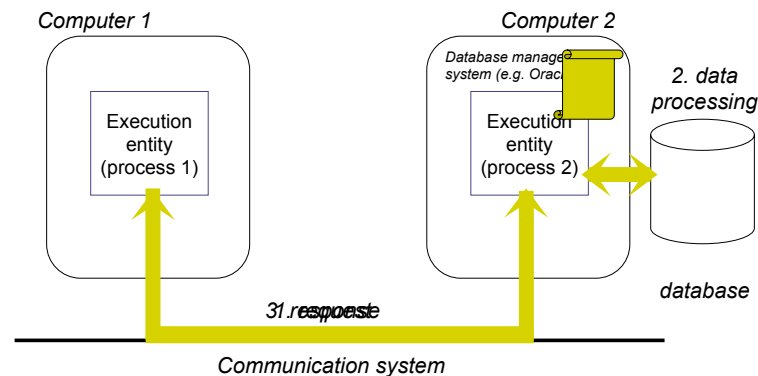
Communication mechanisms in a distributed system

- ◆ Remote procedure call (e.g. a web application)



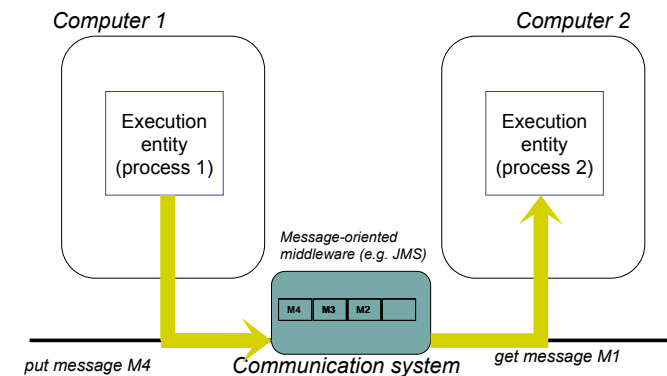
Communication mechanisms in a distributed system

- ◆ Distributed transaction processing (e.g. a database server)



Communication mechanisms in a distributed system

- ◆ Message passing (e.g. a chat system)



Outline

1. What is a distributed system
 - ❖ Communication mechanisms in distributed systems
 - ❖ **Services and interfaces in computing systems**
 - ❖ Client/server architecture
2. What is a middleware
3. References


Services and interfaces in a computing system

- ◆ **Service**
 - ❖ A set of functions: implementing a given behavior, reusable, that are used in a predefined manner
 - Each SW or HW component provides a service
 - A service may be realized in different ways
 - "A service is a contractually defined behavior that can be implemented and provided by any component for use by another component, based solely on the contract", Bieber et. al., Service oriented programming, <http://www.openwings.org/>
- ◆ **Interface**
 - ❖ A service is accessible via one or several interfaces
 - ❖ An interface defines the possible interaction between a service provider and its user


A Real-world Example



Interfaces (1/2)



Interfaces (2/2)

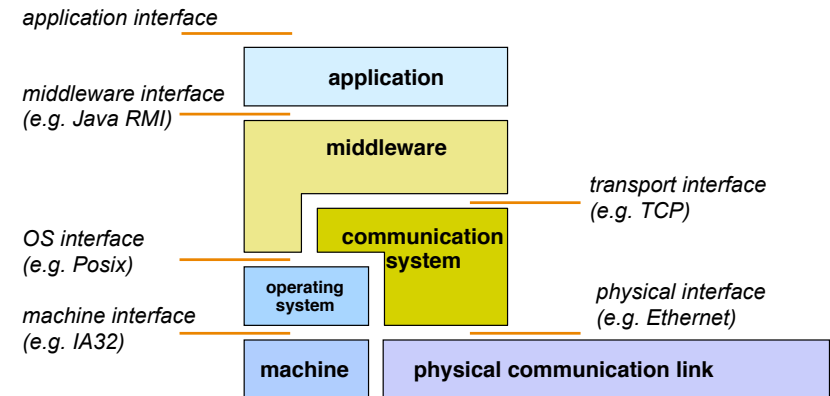


- ◆ A service relies on two interfaces
 - ❖ Required interface (from the service client point of view)
 - ❖ Provided interface (from service provider point of view)
- ◆ Contract
 - ❖ The contract specifies the conformity between the provided and required interfaces
 - ❖ The service client and the service provider are considered as black-boxes; they might be replaced by other implementations as long as the contract is respected
- ◆ The contract may specify aspects that are not related to the interfaces
 - ❖ Non-functional properties related to QoS requirements

Outline

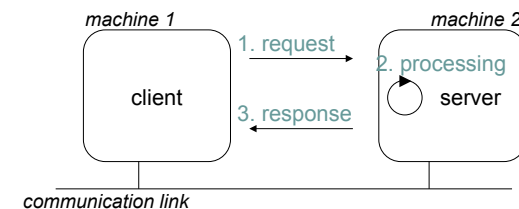
1. What is a distributed system
 - ❖ Communication mechanisms in distributed systems
 - ❖ Services and interfaces in computing systems
 - ❖ **Client/server architecture**
2. What is a middleware
3. References

Examples of important interfaces in computing systems

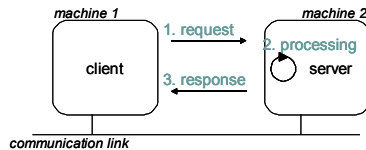


Client/server architecture (1)

- ◆ Definitions
 - ❖ The client/server architecture is a general interaction model
 - ❖ The server provides a service
 - ❖ The client requests that service
 - ❖ The client and the server are usually (but not necessarily) hosted by two distinct machines
 - ❖ Examples of protocols based on the client/server architecture: RPC, Java RMI, Web Services, etc.



Client/server architecture (2)



- ◆ Request message:
 - ❖ Sent by the client to the server
 - ❖ Specifies the requested service (a server may provide several services)
 - ❖ Contains parameters of the requested service
- ◆ Response message:
 - ❖ Sent by the server to the client
 - ❖ Results of service execution, or error message
- ◆ Synchronous communication between the server and the client:
 - ❖ When the client sends a request, it waits (it is blocked) until the server replies to its request

Client/server architecture (3)

◆ Advantages of the client/server architecture

- ❖ Structuring
 - Separation between the interface of a service and the implementation of that service
 - Based on this separation, the client and server implementations can be modified as long as the interface is kept unchanged
- ❖ Protection/security
 - The client and server run in different protection domains
- ❖ Resource management
 - A server may be shared by several clients

Client/server architecture (4)


◆ A server shared by several clients

- ❖ The client point of view



- ❖ The server point of view

- Selecting a request among client requests
- Request processing model (sequential or parallel)



Client/server architecture (5)

◆ Request selection (i.e. scheduling) model

- ❖ First, the server selects one of the waiting (i.e. queued) client requests
- ❖ Then, it process the client request and builds its response
- ❖ Before it returns it to the client

◆ Different request selection strategies

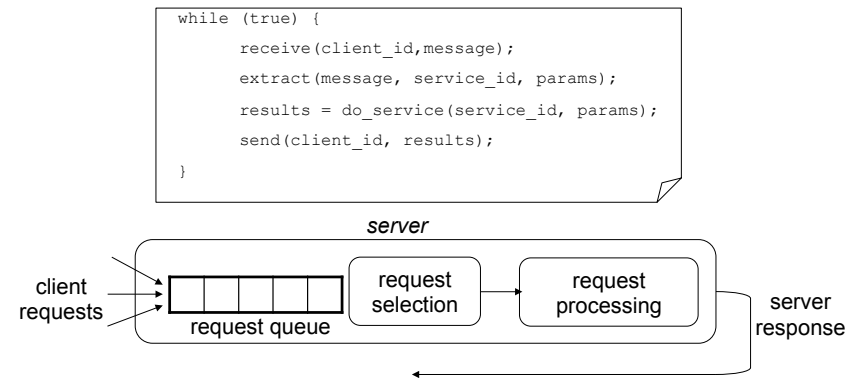
- ❖ First-In First-Out (FIFO)
- ❖ Shortest first
- ❖ Priority-based scheduling

Client/server architecture (6)

- ◆ Request processing model (resource management)
 - ❖ The client and server are executed by two distinct processes (asynchronous call)
 - ❖ The client waits until it receives a response to its request
 - ❖ Several requests may be processed concurrently by the server
 - real parallelism (e.g. multiprocessors, I/O)
 - pseudo-parallelism
 - ❖ Concurrency may take the form of:
 - multiple processes, or
 - multiple threads

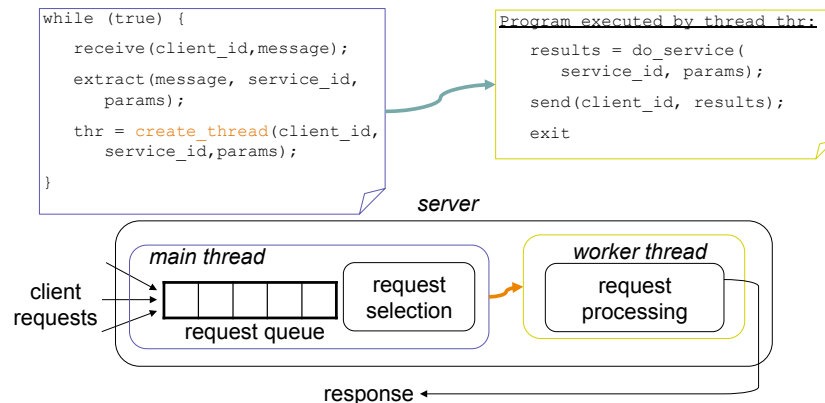
Client/server architecture (7)

- ◆ Server resource management – A unique process



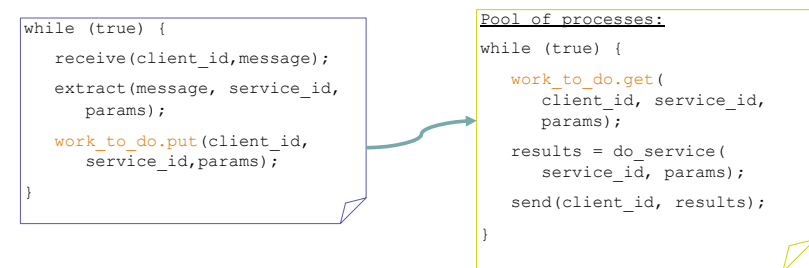
Client/server architecture (8)

- ◆ Server resource management – Multiple processes



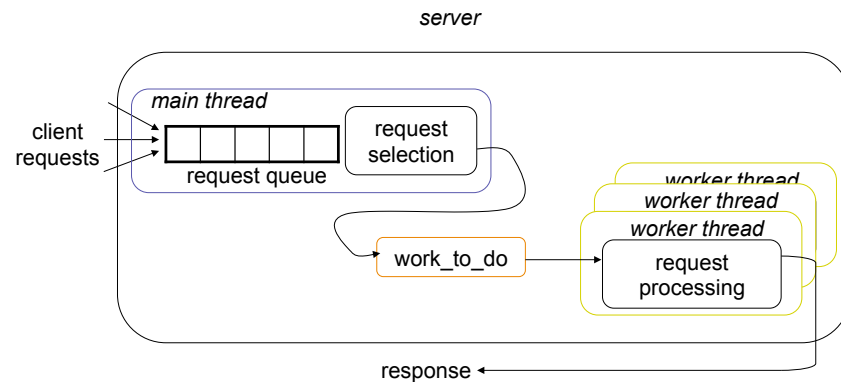
Client/server architecture (9)

- ◆ Server resource management – A pool of processes



Client/server architecture (10)

◆ Server resource management – A pool of processes



Client/server architecture (11)

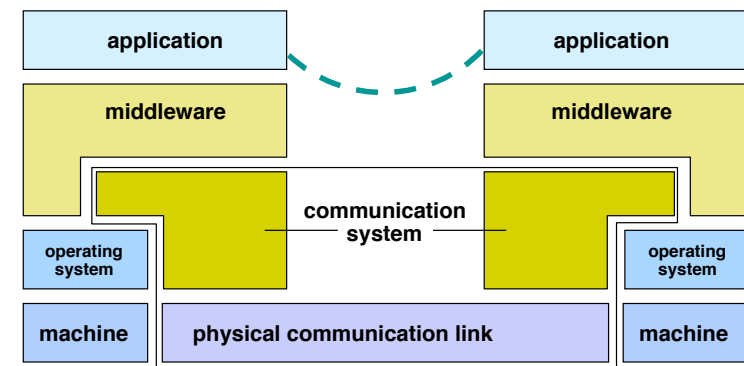
◆ Application of the client/server architecture

- ❖ With low level operations
 - Using functions of the communication system
 - Example: Sockets
 - TCP, connected mode
 - UDP, unconnected mode
- ❖ With high level operations
 - Using a middleware
 - Example: RMI in object-oriented middleware
 - Remote method invocation

Outline

1. What is a distributed system
 - ❖ Communication mechanisms in distributed systems
 - ❖ Services and interfaces in computing systems
 - ❖ Client/server architecture
2. What is a middleware
3. References

What is a middleware



Functions of a middleware


- ◆ A middleware has mainly four functions
 - ❖ Make **distribution as invisible** (transparent) **as possible**
 - ❖ Provide a **homogeneous view** of underlying heterogeneous hardware and software systems
 - ❖ Provide **services of common use** for distributed systems
 - ❖ Provide a **high-level interface** or API (*Applications Programming Interface*) for programming distributed applications
- ◆ Middleware aims at simplifying programming distributed systems
 - ❖ Implementation, evolution and reuse of applications code
 - ❖ Inter-platform portability of applications
 - ❖ Interoperability between heterogeneous applications

References

- ◆ Chris Britton, Peter Bye. *IT Architectures and Middleware: Strategies for Building Large, Integrated Systems (2nd Edition)*. Addison-Wesley, 2004.
- ◆ George Coulouris, Jean Dollimore, Tim Kindberg. *Distributed Systems: Concepts and Design (4th Edition)*. Addison Wesley, 2005.
- ◆ Arno Puder, Kay Römer, Frank Pilhofer. *Distributed Systems Architecture: A Middleware Approach*. Morgan Kaufmann, 2005.
- ◆ Andrew S. Tanenbaum, Maarten van Steen. *Distributed Systems: Principles and Paradigms (2nd Edition)*. Prentice Hall, 2006.
- ◆ This lecture is partly based on lectures given by Sacha Krakowiak, <http://sardes.inrialpes.fr/people/krakowia/>


Examples of middleware solutions

◆ Computation intensive systems



MPI
OpenPBS
Globus
gLite
OpenStack
STORM
...

◆ Information systems



Oracle
MongoDB
pH1
Google Datastore
...