

Impact of memory allocation on the performance of a delegation synchronization algorithm

Presented by **SID-LAKHDAR Riyane**

(M2 MoSIG: ENSIMAG / UJF)

Supervised by **ROPARS Thomas**

(LIG, team ERODS)



MOTIVATION: THE MULTITHREADING LIMITATION




- Theoretical objective of multithreading:

$$Time(\text{multithreaded task}) = \frac{Time(\text{Sequential task})}{\text{Number of thread}}$$

- Objective is still way of the mark

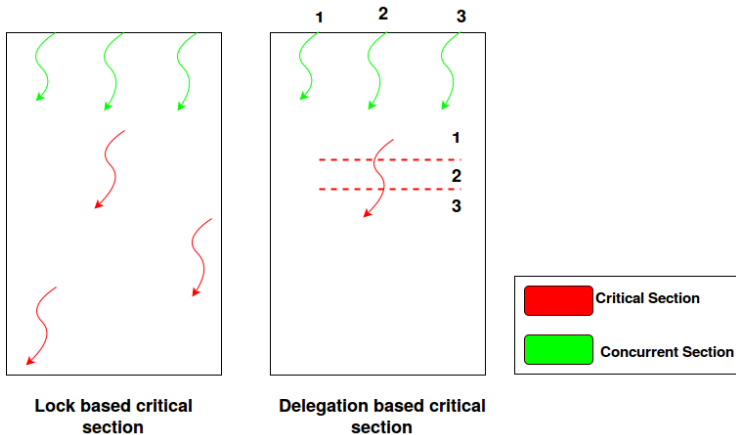
MOTIVATION: COMPLEX PROCESSOR \rightarrow COMPLEX ALLOCATOR



- Local access: ~5 cycles
- Intercore access: ~40 cycles
- Intersocket access: ~150 cycles

- **Memory allocator** is an important cause of this limitation

MOTIVATION: DELEGATION APPROACH



OUR CONTRIBUTION

Prove the impact of dynamic memory allocation on

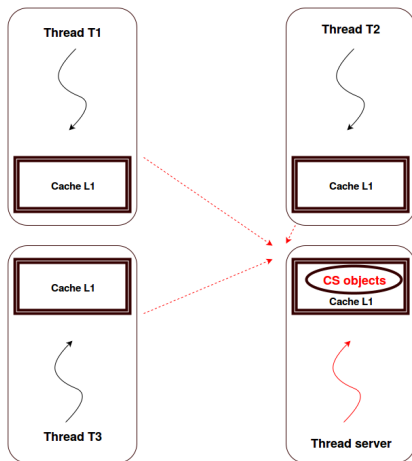
- ▶ The absolute performances of a custom multithreaded queue
- ▶ The relative performance of two custom delegation algorithms

THE DELEGATION APPROACH FOR MUTUAL EXCLUSION

MEMORY ALLOCATION: CHALLENGES AND PROPOSED
SOLUTIONS

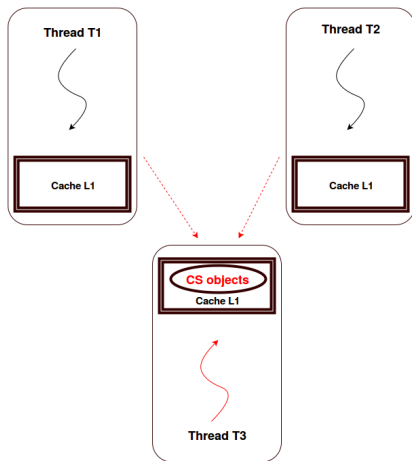
EXPERIMENTAL EVALUATIONS

THE DELEGATION APPROACH




- **Improve cache locality:** memory accesses of the CS are mostly locals

THE COMBINER ALGORITHM



THE EXPECTED OPTIMIZATION: STREAMING STORE



LIMIT OF THE EVALUATION: THE ALLOCATOR

Used allocator:

- ▶ Large thread-local buffer used for all the allocations
- ▶ No returned memory

Unrealistic allocator:


- ▶ Outperforms the general purpose allocators
- ▶ Cannot scale (memory footprint)

THE DELEGATION APPROACH FOR MUTUAL EXCLUSION

**MEMORY ALLOCATION: CHALLENGES AND PROPOSED
SOLUTIONS**


EXPERIMENTAL EVALUATIONS

MEMORY ALLOCATOR: A BOTTLENECK




- ▶ Sequential memory allocation
- ▶ Inefficient due to contention

CORE LOCAL BUFFER



- ▶ Improved allocation time
- ▶ Buffer made of contiguous addresses → low page-fault probability

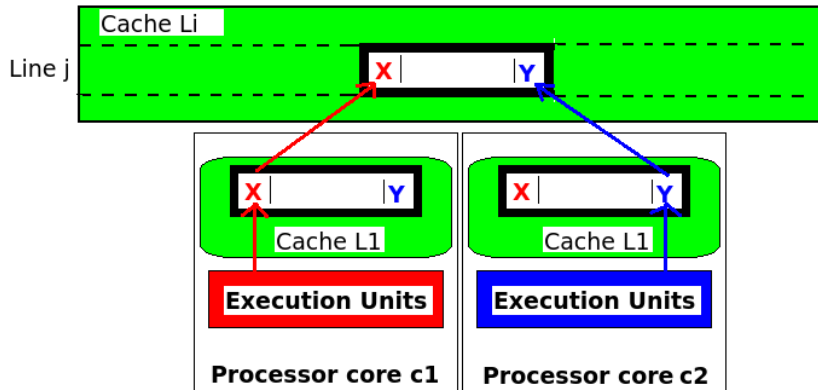
FURTHER CONTENTION LIGHTENING



One stack per core:

- ▶ No synchronization between cores
- ▶ Lock free concurrent data structure

SOFTWARE ISSUES: FALSE SHARING



THE SPAN STRUCTURE

- ▶ **Span size is a multiple of a cache line size:** no false sharing
- ▶ **Simplified allocation algorithm**
- ▶ **Extra memory aligned to page size:** swapped out of the RAM

THE DELEGATION APPROACH FOR MUTUAL EXCLUSION

MEMORY ALLOCATION: CHALLENGES AND PROPOSED
SOLUTIONS

EXPERIMENTAL EVALUATIONS

DESIGNED TEST ENVIRONMENT

Two main rules: Avoid overhead linked to **scheduling**

- ▶ **No more threads than cores**
- ▶ **Pin threads on independent cores**

Environments

- ▶ A 20 cores **Intel Xeon E5-2620** with a **Debian** kernel 3.16.7
- ▶ A 24 cores **AMD Opteron 6164** with a **Debian** kernel 3.16.7

TESTBED

Tested algorithms:


- ▶ Compare contention on a **Michael-Scott** queue
- ▶ Compare the **combiner** and the **server based** algorithms.

Considered memory allocators:


- ▶ **Custom allocator** (used in previous delegation evaluation)
- ▶ **ptmalloc** (C standard library)
- ▶ **hoard**
- ▶ **jemalloc**
- ▶ **scalloc**
- ▶ **supermalloc**
- ▶ **tcmalloc**

ALLOCATOR COMPARISON (WITH THE COMBINER ALGORITHM)

Compare the throughput of the **allocators**: The higher the better



(f) Intel Xeon E5-2620




(g) AMD Opteron 6164


► tcmalloc outperforms the other general purpose allocators

DELEGATION ALGORITHMS COMPARISON

Compare the throughput of the **delegation algorithms**: The higher the better



(h) Intel Xeon E5-2620



(i) AMD Opteron 6164

- With TCMalloc, the relative performance of the algorithms change on AMD

RECAP

During this study, we have:

- ▶ Highlighted the impact of memory allocation on multithreaded algorithms
- ▶ Informed the performance property of two delegation algorithms

Open issues: understand the reasons of the noticed behavior

- ▶ The allocation design of tcmalloc
- ▶ The processor allocation policy

Impact of memory allocation on the performance of a delegation synchronization algorithm

Presented by **SID-LAKHDAR Riyane**

(M2 MoSIG: ENSIMAG / UJF)

Directed by **ROPARS Thomas**

(LIG, team ERODS)



June 21, 2016