

Network lab #4

Transport Protocols : UDP / TCP

Ensimag – Phelma – Master CSE

1 Foreword

During this lab course you will have to go through a series of configurations and observations that will help you understand network principles.



0 – This icon indicates that you must configure the network. Below is the first configuration to do before going further! Before starting any configuration, you must isolate your workstations from the university network they are connected to through interface **bge0** using the following commands:

```
# ifconfig bge0 down
# ifconfig bge0 delete
```



0 – This icon indicates that you have questions to answer. There is a blank space on this sheet to write down your answers.

Making experiments in networking requires great attention and thoroughness. Be sure to stop the commands that are running once you have finished your observations. Try to use several terminals and start only one command in each of them so that you really control what is going on. Have **tcpdump** running to check what is going on in the network, and for example that all the commands generating traffic were stopped correctly before proceeding to the next experiment.

2 Introduction

During this lab, the four stations must belong to the same network.

3



1 – Use two stations, execute **socklab** and say that you want to work with UDP:

```
# socklab udp
```

Create an UDP socket on each station.

```
socklab-udp> sock
```

These two sockets will be used to exchange data between two stations.

This is equivalent to: **sock udp**, creation of a UDP *socket* and **bind**, assignment of an IP address and a port.

Write down this port number: it identifies the socket in your machine in a unique manner.



2 – On one of these two stations, ask for the emission of a data packet to the other station. Of course, you must specify the name and port number of the destination:

```
socklab-udp> sendto <Socket Id> <hostname> <port #>
```

Specify the string that must be transmitted to the destination.



3 – What is the goal of the socket identifier?



4 – On the second station, ask to receive a packet on the socket you created, and specify the number of bytes you want to read:

```
socklab-udp> recvfrom <Socket Id> <byte #>
```

Using **ethereal**, analyze the packet(s) generated by this emission.

TIP: In the **capture** menu select **options** and tick the capture in real-time for a live analysis of the network. Moreover, to avoid **ethereal** filtering information go into **Preferences** → **protocols** → TCP and uncheck everything.



5 – Specify the role of each field in the UDP header. What are the information transmitted by UDP to IP (data + service parameters)?

A UDP packet has the following structure :

| | | |
|---------------------|-----------------------|----|
| 0 | 16 | 31 |
| Source port (16) | Destination port (16) | |
| Message length (16) | Checksum (16) | |



6 – Retry the previous operations (emission and reception) by swapping the role of the two stations (one socket can be used for both reception and emission). You can try the following variants:

- ask to receive the data before the emission;
- send several packets before trying to read them;
- ask for one emission from both sides: Execute the emissions of both extremities, and try to receive them;
- observe what is happening when you try to read less/more data than what is emitted;
- send a data packet to a station disconnected from the network;
- Send several large packets so that the reception buffer is full: e.g., send 5 packets of 9000 bytes, and Try to receive them (with parameter **#9000** to send packets of 9000 bytes).



7 – Try to sum up the behavior of UDP.

4 TCP

Reminder : TCP is a connection oriented protocol, and guaranteed *byte flows* without errors.

TCP headers are :

| | | | | |
|-----------------------------|--------------|------------------|-----------------------|----|
| 0 | 4 | 10 | 16 | 31 |
| Source port (16) | | | Destination port (16) | |
| Sequence number (32) | | | | |
| Acknowledgement number (32) | | | | |
| Hlen (4) | Reserved (6) | Control bits (6) | Window (16) | |
| Checksum (16) | | | Urgent pointer (16) | |
| Options (8.n) | | Padding | | |

The field **control bits** is divided in 6 *flags* : URG, ACK, PSH, RST, SYN, FIN.

Fields signification:

source port and **destination port** have a similar meaning as in UDP;

sequence number and **ack number** are used for the sequencing and error control (the flag ACK indicates if the field **ack number** contains a valid value);

hlen specifies the header length in words of 4 bytes (the TCP header length could vary : one or several optional fields of 4 bytes can be present);

checksum is the same as in UDP;

urgent pointer is used to carry urgent data (the flag URG specifies if the field **urgent pointer** is a valid value);

the flags SYN and FIN are used to open and close virtual connections ; RST is used to reinitialize the virtual connections in uncertain state (SYN duplicated or failed);

window is used for flow control.

4.1 Connection establishment



8 – In two stations, execute **socklab** with TCP :

```
# socklab tcp
```

On the first machine, create a *passive* socket. Write down the port number and put it in listen mode.

```
socklab-tcp> passive
```

It is equivalent to : **sock tcp**, socket creation, **bind**, IP address and port number assignment, **listen**, waits on this socket.



9 – Put the socket in a mode where it waits to accept a connection:

```
socklab-tcp> accept
```

On the second machine, create an *active* socket and connect it to the first machine (with the *passive* socket). You can use :

```
socklab-tcp> connect <hostname> <port #>
```

It is equivalent to : **sock tcp**, socket creation, **bind**, IP address and port number assignment and **connect**, connection request.



10 – Why the first socket is called *passive*, and the second one *active* ? What are the role of these two sockets in the connection opening?



11 – Analyze the packets for this connection.



12 – Extract the steps of this connection, how they are organized, messages exchanged, with a temporal diagram



13 – Explain what happens when the socket is in the *accept* mode. Try to do it before and after the *connect*. Take a look on the socket list (status command).



14 – What information is exchanged ? What is the goal of the flag SYN ?



15 – Open several connections to the same port destination.



16 – What is the real identifier of one connection, how TCP associated the received messages to one connection?

4.2 Connection termination



17 – After opening a connection between two hosts, close it:
`socklab-tcp> close`



18 – Analyze the packets and their direction. Find the steps of this closing. What is the goal of the flag FIN ?



19 – Sum up the messages exchanges and specify the field values for this step.



20 – Summary: Write the state machine for TCP active and passive sockets (cf. Figure 1).




Figure 1: State machine example

4.3 Sequencing and error control



21 – Establish a TCP connection between the two stations and exchange some data:
`socklab-tcp> read <byte #>`
 to read bytes on the new socket, and
`socklab-tcp> write`
 to send messages (instead of a normal message, you can use the notation `#n` to send a message of n bytes).



22 – Analyze this exchange. Explain the goal of the fields **sequence number** and **ack number** in TCP packets.



23 – Give an intuitive explanation of the update rules for the value of these two fields. Distinguish the case in which we receive a packet from the case in which we emit one packet.



24 – Unplug the cable between the two stations (disconnected stations) and redo the operation. What happens? Reconnect the stations, check the traffic.



25 – Explain what happens in this kind of situation (segment loss), compare with UDP.

4.4 Flow control



26 – Open a TCP connection. Send a massive amount of data on this connection (70'000 bytes).



27 – What is happening? Why?



28 – Execute several successive **read** at the receiver (with 1 byte then 100 and then 5'000).



29 – Analyze the last exchanged packets. Take a closer look on the field **window**. Explain how the control flow works.

5 IP fragmentation

We will now take a look on the IP fragmentation mechanism.

IP fragmentation is used when an IP packet is too large to be transmitted as is through the link layer

(an Ethernet link cannot deal with packets longer than 1518 bytes : Ethernet header and CRC included). Several small fragments must be created.

When a packet is fragmented, it must be recombined at the receiver (n the IP layer). If the destination does not support re-assembly, fragmentation must be forbidden (with the flag DF, *Don't Fragment*). When this flag is set, fragmentation is forbidden. The delivery could be impossible. We will now analyze this fragmentation.



30 – Execute **ethereal** on the monitoring station. On another station, send a UDP segment of 4'600 bytes to a socket that was previously created.



31 – Analyze exchanges. In particular, study the role of **total length**, **fragment offset**, and the flag MF, *More Fragments*.



32 – Repeat this manipulation with longer packets (if necessary). Repeat the same operations with TCP.



33 – Conclusions?



34 – Verify the conclusions about **sequence number** and **ack number** fields of the TCP header previously described.